

SQL (Structured Query Language) Practice

Employees Table

| EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID |
|------------|-----------|----------|--------------|--------|------------|-----------|
| 1 | Alice | Johnson | 101 | 60000 | 2018-01-15 | 3 |
| 2 | Bob | Smith | 102 | 75000 | 2017-05-20 | 3 |
| 3 | Charlie | Brown | 101 | 90000 | 2015-09-30 | NULL |
| 4 | David | Williams | 103 | 55000 | 2019-07-11 | 3 |
| 5 | Eva | Davis | 102 | 65000 | 2020-03-25 | 2 |

Orders Table

| OrderID | EmployeeID | ProductID | Quantity | OrderDate |
|---------|------------|-----------|----------|------------|
| 1001 | 1 | 201 | 10 | 2022-01-15 |
| 1002 | 2 | 202 | 5 | 2022-01-16 |
| 1003 | 3 | 203 | 20 | 2022-01-17 |
| 1004 | 4 | 202 | 15 | 2022-01-18 |
| 1005 | 5 | 204 | 25 | 2022-01-19 |

Products Table

| ProductID | ProductName | Price | Category |
|-----------|--------------|-------|-------------|
| 201 | Laptop | 1200 | Electronics |
| 202 | Smartphone | 800 | Electronics |
| 203 | Office Chair | 150 | Furniture |
| 204 | Desk | 300 | Furniture |
| 205 | Monitor | 200 | Electronics |

SQL Queries Interview Questions

Q1. Write a query to display all records from the Employees table.

Answer:

SELECT * FROM Employees;Copy Code

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

| | | | | | | |
|---|---------|----------|----|--------|------------|------|
| 1 | Alice | Johnson | 10 | 160000 | 2018-01-15 | 32 |
| 2 | Bob | Smith | 10 | 275000 | 2017-05-20 | 33 |
| 3 | Charlie | Brown | 10 | 190000 | 2015-09-30 | NULL |
| 4 | David | Williams | 10 | 355000 | 2019-07-11 | 35 |
| 5 | Eva | Davis | 10 | 265000 | 2020-03-25 | 32 |

Copy Code

Q2. Fetch only the FirstName and LastName of employees.

Answer:

SELECT FirstName, LastName FROM Employees;

FirstName | LastName

| | |
|---------|----------|
| Alice | Johnson |
| Bob | Smith |
| Charlie | Brown |
| David | Williams |
| Eva | Davis |

Copy Code

Q3. Retrieve the unique department IDs from the Employees table.

Answer:

SELECT DISTINCT DepartmentID FROM Employees;

DepartmentID

10

Copy Code

Q4. Fetch employees with a salary greater than 60,000.

Answer:

```
SELECT * FROM Employees WHERE Salary > 60000;
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

| | | | | | | |
|---|---------|----------|----|--------|------------|------|
| 1 | Alice | Johnson | 10 | 160000 | 2018-01-15 | 32 |
| 2 | Bob | Smith | 10 | 275000 | 2017-05-20 | 33 |
| 3 | Charlie | Brown | 10 | 190000 | 2015-09-30 | NULL |
| 4 | David | Williams | 10 | 355000 | 2019-07-11 | 35 |
| 5 | Eva | Davis | 10 | 265000 | 2020-03-25 | 32 |

Copy Code

Q5. Write a query to display all orders placed on or after January 17, 2022.

Answer:

```
SELECT * FROM Orders WHERE OrderDate >= '2022-01-17';
```

OrderID | EmployeeID | ProductID | Quantity | OrderDate

| | | | | |
|------|---|---|---|------------|
| 1022 | 2 | 1 | 2 | 2022-01-16 |
| 1023 | 3 | 3 | 3 | 2022-01-17 |
| 1024 | 4 | 2 | 5 | 2022-01-18 |
| 1025 | 5 | 4 | 5 | 2022-01-19 |

Copy Code

Q6. Retrieve all products with a price less than 300.

Answer:

SELECT * FROM Products WHERE Price < 300;

ProductID | ProductName | Price | Category

203 | Office Chair | 150 | Furniture

204 | Desk | 300 | Furniture

205 | Monitor | 200 | Electronics

Copy Code

Q7. Find the total number of orders in the Orders table.

Answer:

SELECT COUNT(*) AS TotalOrders FROM Orders;

TotalOrders

5

Copy Code

Q8. Fetch the details of the product named 'Laptop'.

Answer:

SELECT * FROM Products WHERE ProductName = 'Laptop';

ProductID | ProductName | Price | Category

201 | Laptop | 1200 | Electronics

Copy Code

Q9. Write a query to sort employees by their HireDate in ascending order.

Answer:

SELECT * FROM Employees ORDER BY HireDate ASC;

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

| | | | | | | |
|---|---------|----------|----|--------|------------|------|
| 3 | Charlie | Brown | 10 | 190000 | 2015-09-30 | NULL |
| 2 | Bob | Smith | 10 | 275000 | 2017-05-20 | 33 |
| 1 | Alice | Johnson | 10 | 160000 | 2018-01-15 | 32 |
| 4 | David | Williams | 10 | 355000 | 2019-07-11 | 35 |
| 5 | Eva | Davis | 10 | 265000 | 2020-03-25 | 32 |

Copy Code

Q10. Retrieve the maximum price of products in the Electronics category.

Answer:

```
SELECT MAX(Price) AS MaxPrice FROM Products WHERE Category = 'Electronics';
```

MaxPrice

1200Copy Code

Q11. Write a query to join Employees and Orders tables to fetch employee names along with their orders.

Answer:

```
SELECT e.FirstName, e.LastName, o.OrderID, o.OrderDate
FROM Employees e
JOIN Orders o ON e.EmployeeID = o.EmployeeID;
```

FirstName | LastName | OrderID | OrderDate

Alice | Johnson | 1022 | 2022-01-16

Bob | Smith | 1023 | 2022-01-17

Charlie | Brown | 1024 | 2022-01-18

David | Williams | 1025 | 2022-01-19

Copy Code

Q12. Calculate the total salary by department.

Answer:

```
SELECT DepartmentID, SUM(Salary) AS TotalSalary  
FROM Employees  
GROUP BY DepartmentID;
```

DepartmentID | TotalSalary

10 | 1355000

Copy Code

Q13. Find the employees who do not have a manager.

Answer:

```
SELECT * FROM Employees WHERE ManagerID IS NULL;
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

3 | Charlie | Brown | 10 | 190000 | 2015-09-30 | NULL

Copy Code

Q14. Write a query to display the average product price for each category.

Answer:

```
SELECT Category, AVG(Price) AS AvgPrice  
FROM Products  
GROUP BY Category;
```

Category | AvgPrice

Electronics | 800

Furniture | 216.67

Copy Code

Q15. Fetch the details of the top 3 highest-paid employees.

Answer:

```
SELECT * FROM Employees  
ORDER BY Salary DESC  
LIMIT 3;
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

| | | | | | | |
|---|-------|----------|----|--------|------------|----|
| 4 | David | Williams | 10 | 355000 | 2019-07-11 | 35 |
| 2 | Bob | Smith | 10 | 275000 | 2017-05-20 | 33 |
| 5 | Eva | Davis | 10 | 265000 | 2020-03-25 | 32 |

Copy Code

Q16. Retrieve the order details along with the product name.

Answer:

```
SELECT o.OrderID, o.Quantity, p.ProductName, p.Price  
FROM Orders o  
JOIN Products p ON o.ProductID = p.ProductID;
```

OrderID | Quantity | ProductName | Price

1022 | 2 | Laptop | 1200

1023 | 3 | Office Chair | 150

1024 | 5 | Smartphone | 800

1025 | 5 | Desk | 300

Copy Code

Q17. Find the total quantity of products ordered for each product.

Answer:

```
SELECT ProductID, SUM(Quantity) AS TotalQuantity
```

```
FROM Orders
```

```
GROUP BY ProductID;
```

ProductID | TotalQuantity

1 | 2

2 | 8

3 | 3

4 | 5

Copy Code

Q18. Write a query to update the price of all Furniture category products by 10%.

Answer:

```
UPDATE Products
```

```
SET Price = Price * 1.10
```

```
WHERE Category = 'Furniture'; Copy Code
```

Q19. Delete all orders placed before January 17, 2022.

Answer:

```
DELETE FROM Orders WHERE OrderDate < '2022-01-17'; Copy Code
```

Q20. Fetch employees whose first name starts with 'A'.

Answer:

```
SELECT * FROM Employees WHERE FirstName LIKE 'A%';
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

1 | Alice | Johnson | 10 | 160000 | 2018-01-15 | 32

Copy Code

Q21. Retrieve the number of employees hired each year.

Answer:

```
SELECT YEAR(HireDate) AS HireYear, COUNT(*) AS EmployeesHired  
FROM Employees  
GROUP BY YEAR(HireDate);
```

HireYear | EmployeesHired

2015 | 1

2017 | 1

2018 | 1

2019 | 1

2020 | 1

Copy Code

Q22. Write a query to fetch employees earning more than the average salary.

Answer:

```
SELECT * FROM Employees  
WHERE Salary > (SELECT AVG(Salary) FROM Employees);
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

2 | Bob | Smith | 10 | 275000 | 2017-05-20 | 33

| | | | | | | |
|---|-------|----------|----|--------|------------|----|
| 4 | David | Williams | 10 | 355000 | 2019-07-11 | 35 |
| 5 | Eva | Davis | 10 | 265000 | 2020-03-25 | 32 |

Copy Code

Q23. Display the top 3 products with the highest total quantity sold.

Answer:

```
SELECT p.ProductName, SUM(o.Quantity) AS TotalQuantity  
FROM Orders o  
JOIN Products p ON o.ProductID = p.ProductID  
GROUP BY p.ProductName  
ORDER BY TotalQuantity DESC  
LIMIT 3;
```

ProductName | TotalQuantity

Smartphone | 8

Desk | 5

Office Chair | 3

Copy Code

Q24. Retrieve the employees who have not placed any orders.

Answer:

```
SELECT * FROM Employees  
WHERE EmployeeID NOT IN (SELECT DISTINCT EmployeeID FROM Orders);
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

| | | | | | | |
|---|---------|-------|----|--------|------------|------|
| 3 | Charlie | Brown | 10 | 190000 | 2015-09-30 | NULL |
|---|---------|-------|----|--------|------------|------|

Copy Code

Q25. Write a query to fetch the most recently hired employee.

Answer:

```
SELECT * FROM Employees
ORDER BY HireDate DESC
LIMIT 1;
```

EmployeeID | FirstName | LastName | DepartmentID | Salary | HireDate | ManagerID

5 | Eva | Davis | 10 | 265000 | 2020-03-25 | 32

Copy Code

Q26. Display all employees along with the total number of orders they've handled.

Answer:

```
SELECT e.EmployeeID, e.FirstName, COUNT(o.OrderID) AS TotalOrders
FROM Employees e
LEFT JOIN Orders o ON e.EmployeeID = o.EmployeeID
GROUP BY e.EmployeeID, e.FirstName;Copy Code
```

| EmployeeID | FirstName | TotalOrders |
|------------|-----------|-------------|
| 1 | Alice | 2 |
| 2 | Bob | 2 |
| 3 | Charlie | 1 |
| 4 | David | 1 |
| 5 | Eva | 0 |

Q27. Fetch product details for which total sales exceed \$10,000.

Answer:

```
SELECT p.ProductName, SUM(o.Quantity * p.Price) AS TotalSales
FROM Orders o
JOIN Products p ON o.ProductID = p.ProductID
```

GROUP BY p.ProductName

HAVING TotalSales > 10000;Copy Code

| ProductName | TotalSales |
|-------------|------------|
| Laptop | 24000 |

Q28. Find employees who joined the company in the same year as their manager.

Answer:

```
SELECT e.FirstName AS EmployeeName, m.FirstName AS ManagerName
```

```
FROM Employees e
```

```
JOIN Employees m ON e.ManagerID = m.EmployeeID
```

```
WHERE YEAR(e.HireDate) = YEAR(m.HireDate);Copy Code
```

| EmployeeName | ManagerName |
|--------------|-------------|
| Alice | Bob |

Q29. Retrieve the employee names with the highest salary in each department.

Answer:

```
SELECT DepartmentID, FirstName, LastName, Salary
```

```
FROM Employees
```

```
WHERE (DepartmentID, Salary) IN (
```

```
    SELECT DepartmentID, MAX(Salary)
```

```
    FROM Employees
```

```
    GROUP BY DepartmentID
```

```
);Copy Code
```

| DepartmentID | FirstName | LastName | Salary |
|--------------|-----------|----------|--------|
| 1 | Alice | Johnson | 160000 |
| 2 | Bob | Smith | 75000 |
| 3 | David | Williams | 55000 |

Q30. Write a query to fetch the total revenue generated by each employee.

Answer:

```

SELECT e.FirstName, e.LastName, SUM(o.Quantity * p.Price) AS TotalRevenue
FROM Employees e
JOIN Orders o ON e.EmployeeID = o.EmployeeID
JOIN Products p ON o.ProductID = p.ProductID
GROUP BY e.EmployeeID, e.FirstName, e.LastName;Copy Code

```

| FirstName | LastName | TotalRevenue |
|-----------|----------|--------------|
| Alice | Johnson | 32000 |
| Bob | Smith | 63000 |
| Charlie | Brown | 45000 |
| David | Williams | 30000 |
| Eva | Davis | 0 |

Q31. Write a query to fetch employees earning more than their manager.

Answer:

```

SELECT e.FirstName AS EmployeeName, m.FirstName AS ManagerName
FROM Employees e
JOIN Employees m ON e.ManagerID = m.EmployeeID
WHERE e.Salary > m.Salary;Copy Code

```

| EmployeeName | ManagerName |
|--------------|-------------|
| Alice | Bob |

Q32. Retrieve the second highest salary from the Employees table.

Answer:

```

SELECT MAX(Salary) AS SecondHighestSalary
FROM Employees
WHERE Salary < (SELECT MAX(Salary) FROM Employees);

```

SecondHighestSalary

75000Copy Code

Q33. List the departments with no employees assigned.

Answer:

```
SELECT * FROM Departments
```

```
WHERE DepartmentID NOT IN (SELECT DISTINCT DepartmentID FROM Employees);Copy  
Code
```

| DepartmentID | DepartmentName |
|--------------|----------------|
| 4 | Marketing |

Q34. Write a query to create a view showing employee names and their department names.

Answer:

```
CREATE VIEW EmployeeDepartmentView AS
```

```
SELECT e.FirstName, e.LastName, d.DepartmentName
```

```
FROM Employees e
```

```
JOIN Departments d ON e.DepartmentID = d.DepartmentID;Copy Code
```

| FirstName | LastName | DepartmentName |
|-----------|----------|----------------|
| Alice | Johnson | IT |
| Bob | Smith | Sales |
| Charlie | Brown | IT |
| David | Williams | HR |
| Eva | Davis | Sales |

Q35. Fetch the names of employees who have placed more than 10 orders.

Answer:

```
SELECT e.FirstName, e.LastName
```

```
FROM Employees e
```

```
JOIN Orders o ON e.EmployeeID = o.EmployeeID
```

```
GROUP BY e.EmployeeID, e.FirstName, e.LastName
```

```
HAVING COUNT(o.OrderID) > 10;Copy Code
```

| FirstName | LastName |
|-----------|----------|
| Alice | Johnson |
| Bob | Smith |

Q36. Write a query to rank employees based on their salary within each department.

Answer:

```
SELECT EmployeeID, FirstName, DepartmentID, Salary,
       RANK() OVER (PARTITION BY DepartmentID ORDER BY Salary DESC) AS Rank
FROM Employees;Copy Code
```

| EmployeeID | FirstName | DepartmentID | Salary | Rank |
|------------|-----------|--------------|--------|------|
| 1 | Alice | 1 | 160000 | 1 |
| 3 | Charlie | 1 | 190000 | 2 |
| 2 | Bob | 2 | 75000 | 1 |
| 4 | David | 3 | 55000 | 1 |
| 5 | Eva | 2 | 65000 | 2 |

Q37. Retrieve the cumulative sales for each product.

Answer:

```
SELECT ProductID, ProductName,
       SUM(SUM(Quantity * Price)) OVER (ORDER BY ProductID) AS CumulativeSales
FROM Products p
JOIN Orders o ON p.ProductID = o.ProductID
GROUP BY ProductID, ProductName;Copy Code
```

| ProductID | ProductName | CumulativeSales |
|-----------|--------------|-----------------|
| 201 | Laptop | 24000 |
| 202 | Smartphone | 32000 |
| 203 | Office Chair | 1500 |
| 204 | Desk | 3000 |

| ProductID | ProductName | CumulativeSales |
|-----------|-------------|-----------------|
| 205 | Monitor | 1500 |

Q38. Identify the department with the highest total salary expenditure.

Answer:

```
SELECT DepartmentID, SUM(Salary) AS TotalExpenditure
FROM Employees
GROUP BY DepartmentID
ORDER BY TotalExpenditure DESC
LIMIT 1;Copy Code
```

| DepartmentID | TotalExpenditure |
|--------------|------------------|
| 1 | 450000 |

Q39. Write a query to find the percentage contribution of each product to total sales.

Answer:

```
SELECT p.ProductName,
       (SUM(o.Quantity * p.Price) * 100.0 /
        (SELECT SUM(Quantity * Price) FROM Orders o JOIN Products p ON o.ProductID =
         p.ProductID)) AS ContributionPercentage
FROM Orders o
JOIN Products p ON o.ProductID = p.ProductID
GROUP BY p.ProductName;Copy Code
```

| ProductName | ContributionPercentage |
|--------------|------------------------|
| Laptop | 48.00 |
| Smartphone | 32.00 |
| Office Chair | 4.00 |
| Desk | 8.00 |
| Monitor | 8.00 |

Q40. Find employees who have the same manager and earn more than \$70,000.

Answer:

```

SELECT *
FROM Employees e1
WHERE ManagerID IS NOT NULL
AND Salary > 70000
AND ManagerID IN (
    SELECT ManagerID FROM Employees e2 WHERE e1.ManagerID = e2.ManagerID
);Copy Code

```

| EmployeeID | FirstName | LastName | Salary | ManagerID |
|------------|-----------|----------|--------|-----------|
| 1 | Alice | Johnson | 160000 | 32 |
| 2 | Bob | Smith | 75000 | 32 |

Q41. Write a query to detect duplicate rows in the Orders table.

Answer:

```

SELECT EmployeeID, ProductID, OrderDate, COUNT(*) AS DuplicateCount
FROM Orders
GROUP BY EmployeeID, ProductID, OrderDate
HAVING COUNT(*) > 1;Copy Code

```

| EmployeeID | ProductID | OrderDate | DuplicateCount |
|------------|-----------|------------|----------------|
| 1 | 201 | 2022-01-15 | 2 |

Q42. Fetch the details of orders placed on the same day by multiple employees.

Answer:

```

SELECT OrderDate, COUNT(DISTINCT EmployeeID) AS EmployeeCount
FROM Orders
GROUP BY OrderDate
HAVING EmployeeCount > 1;Copy Code

```

| OrderDate | EmployeeCount |
|------------|---------------|
| 2022-01-15 | 2 |
| 2022-01-16 | 2 |
| 2022-01-17 | 1 |

Q43. Create a stored procedure to update product prices based on category.

Answer:

DELIMITER \$\$

```
CREATE PROCEDURE UpdatePriceByCategory(IN category_name VARCHAR(50), IN
price_factor DECIMAL(5, 2))
```

BEGIN

 UPDATE Products

 SET Price = Price * price_factor

 WHERE Category = category_name;

END\$\$

DELIMITER ;Copy Code

Q44. Write a query to calculate the lead and lag in order dates for each employee.

Answer:

SELECT EmployeeID, OrderID, OrderDate,

```
    LAG(OrderDate) OVER (PARTITION BY EmployeeID ORDER BY OrderDate) AS
PreviousOrderDate,
```

```
    LEAD(OrderDate) OVER (PARTITION BY EmployeeID ORDER BY OrderDate) AS
NextOrderDate
```

FROM Orders;Copy Code

| EmployeeID | OrderID | OrderDate | PreviousOrderDate | NextOrderDate |
|------------|---------|------------|-------------------|---------------|
| 1 | 1 | 2022-01-15 | NULL | 2022-01-16 |
| 2 | 2 | 2022-01-16 | 2022-01-15 | 2022-01-17 |
| 3 | 3 | 2022-01-17 | NULL | NULL |

Q45. Identify the products that have not been ordered.

Answer:

```
SELECT * FROM Products
```

WHERE ProductID NOT IN (SELECT DISTINCT ProductID FROM Orders);Copy Code

| ProductID | ProductName |
|-----------|-------------|
| 204 | Desk |
| 205 | Monitor |

Q46. Write a query to fetch employees whose total order quantity is between 50 and 100.

Answer:

```
SELECT e.FirstName, e.LastName, SUM(o.Quantity) AS TotalQuantity
```

```
FROM Employees e
```

```
JOIN Orders o ON e.EmployeeID = o.EmployeeID
```

```
GROUP BY e.EmployeeID, e.FirstName, e.LastName
```

HAVING TotalQuantity BETWEEN 50 AND 100;Copy Code

| FirstName | LastName | TotalQuantity |
|-----------|----------|---------------|
| Bob | Smith | 60 |

Q47. Fetch the second-highest quantity ordered for each product.

Answer:

```
SELECT ProductID, MAX(Quantity) AS SecondHighestQuantity
```

```
FROM Orders
```

```
WHERE Quantity < (SELECT MAX(Quantity) FROM Orders WHERE Orders.ProductID = ProductID)
```

GROUP BY ProductID;Copy Code

| ProductID | SecondHighestQuantity |
|-----------|-----------------------|
| 201 | 20 |
| 202 | 30 |

| ProductID | SecondHighestQuantity |
|-----------|-----------------------|
| 203 | 10 |

Q48. Find the minimum and maximum order quantities for each employee.

Answer:

```
SELECT EmployeeID, MIN(Quantity) AS MinQuantity, MAX(Quantity) AS MaxQuantity
FROM Orders
GROUP BY EmployeeID;Copy Code
```

| EmployeeID | MinQuantity | MaxQuantity |
|------------|-------------|-------------|
| 1 | 10 | 20 |
| 2 | 20 | 40 |
| 3 | 10 | 10 |

Q49. Write a query to split employee salaries into quartiles.

Answer:

```
SELECT EmployeeID, FirstName, Salary,
       NTILE(4) OVER (ORDER BY Salary) AS SalaryQuartile
FROM Employees;Copy Code
```

| EmployeeID | FirstName | Salary | SalaryQuartile |
|------------|-----------|--------|----------------|
| 1 | Alice | 160000 | 4 |
| 2 | Bob | 75000 | 3 |
| 3 | Charlie | 190000 | 4 |
| 4 | David | 55000 | 2 |
| 5 | Eva | 65000 | 2 |

Q50. Create a temporary table for orders with high revenue (greater than \$5000)

Answer:

```
CREATE TEMPORARY TABLE HighRevenueOrders AS
SELECT o.OrderID, o.Quantity, p.Price, (o.Quantity * p.Price) AS Revenue
```

```
FROM Orders o
JOIN Products p ON o.ProductID = p.ProductID
WHERE (o.Quantity * p.Price) > 5000;Copy Code
```

| OrderID | Quantity | Price | Revenue |
|---------|----------|-------|---------|
| 1 | 10 | 1200 | 12000 |
| 2 | 25 | 800 | 20000 |

able: Person

| Column Name | Type |
|-------------|---------|
| id | int |
| email | varchar |

Here, id is the primary key column for this table. Email represents the email id of the person. For the sake of simplicity, we assume that the emails will not contain uppercase letters. Write an SQL query to report all the duplicate emails. You can return the result table in any order.

Example:

Input: Person table:

| id | email |
|----|-------------|
| 1 | a@gmail.com |
| 2 | c@yahoo.com |
| 3 | a@gmail.com |

Output:

| Email |
|-------------|
| a@gmail.com |

Explanation: a@gmail.com is repeated two times.

- **Approach 1:**

We can first have all the distinct email ids and their respective counts in our result set. For this, we can use the GROUP BY operator to group the tuples by their email id. We will use the COUNT operator to have the total number of a particular email id in the given table. The query for obtaining this resultant set can be written as:

```
select email, count(email) as email_count
from Person
group by email;
```

Now, we query in the above resultant query set to find out all the tuples which have an email id count greater than 1. This can be achieved using the following query:

```
select email from
(
    select email, count(email) as email_count
    from Person
    group by email
)
where email_count > 1;
```

- **Approach 2:**

The HAVING clause, which is significantly simpler and more efficient, is a more popular technique to add a condition to a GROUP BY. So, we can first group the tuples by the email ids and then have a condition to check if their count is greater than 1, only then do we include it in our result set. So we may change the solution above to this one.

```
select email
from Person
group by email
having count(email) > 1;
```

- **Approach 3:**

We can use the concept of joins to solve this problem. We will self-join the Person table with the condition that their email ids should be the same and their ids should be different. Having done this, we just need to count the number of tuples in our resultant set with

distinct email ids. For this, we use the DISTINCT operator. This can be achieved using the following query:

```
SELECT DISTINCT p1.email
FROM Person p1, Person p2
WHERE p1.email = p2.email AND p1.id != p2.id;
```

3. Let us consider the following schema:

Table: Activity

| Column Name | Type |
|-------------|------|
| playerId | int |
| deviceID | int |
| eventDate | date |
| gamesplayed | int |

Example 1:

Input: Activity table:

| PlayerId | deviceID | eventDate | gamesPlayed |
|----------|----------|------------|-------------|
| 1 | 2 | 2021-08-09 | 9 |
| 1 | 2 | 2021-04-07 | 3 |
| 2 | 3 | 2021-06-25 | 1 |
| 3 | 1 | 2021-03-02 | 1 |
| 3 | 4 | 2021-07-03 | 3 |

Output:

| playerId | firstLogin |
|----------|------------|
| 1 | 2021-04-07 |
| 2 | 2021-06-25 |
| 3 | 2021-07-03 |

Explanation:

The player with playerId 1 has two login event dates in the example above. However, because the first login event date is 2021-04-07, we display it. Similarly, the first login event date for the player with playerId 2 is 2021-06-25, and the first login event date for the player with playerId 3 is 2021-07-03.

- **Approach 1:**

We can first group the tuples by their player_id. Now, we want the most initial date when the player logged in to the game. For this, we can use the MIN operator and find the initial date on which the player logged in. The query can be written as follows:

```
select playerId, min(eventDate) as firstLogin from Activity group by playerId
```

- **Approach 2:**

We can partition the tuples by the player_id and order them by their event_id such that all the tuples having the same player_id are grouped together. We then number every tuple in each of the groups starting with the number 1. Now, we just have to display the event_date for the tuple having row number 1. For this, we use the ROW_NUMBER operator. The SQL query for it can be written as follows:

```
SELECT playerId, eventDate AS firstLogin
```

```
FROM
```

```
(
```

```
    SELECT playerId, eventDate, ROW_NUMBER() OVER (PARTITION BY playerId ORDER BY eventDate) AS seq
```

```
    FROM Activity
```

```
) AS t
```

```
WHERE seq = 1
```

- **Approach 3:**

We follow a similar kind of approach as used in Approach 2. But instead of using the ROW_NUMBER operator, we can use the FIRST_VALUE operator to find the first event_date. The SQL query for it can be written as follows:

```
select distinct(playerId),
```

```
    FIRST_VALUE(eventDate) OVER(PARTITION BY playerId ORDER BY eventDate) as firstLogin
    from Activity;
```

You can download a PDF version of Sql Query Interview Questions.

4. Given the following schema:

Table: Customers

| Column Name | Type |
|-------------|---------|
| id | int |
| name | varchar |

The primary key column for this table is id. Each row in the table represents a customer's ID and name.

Table: Orders

| Column Name | Type |
|-------------|------|
| id | int |
| customerId | int |

The primary key column for this table is id. customerId is a foreign key of the ID from the Customers table. The ID of an order and the ID of the customer who placed it are listed in each row of this table. Write an SQL query to report all customers who never order anything. You can return the result table in any order.

Example:

Input: Customers table:

| id | name |
|----|--------|
| 1 | Ram |
| 2 | Sachin |
| 3 | Rajat |
| 4 | Ankit |

Orders table:

| id | customerid |
|-----------|-------------------|
| 1 | 2 |
| 2 | 1 |

Output

| Customers |
|------------------|
| Rajat |
| Ankit |

Explanation: Here, the customers Sachin and Ram have placed an order having order id 1 and 2 respectively. Thus, the customers Rajat and Ankit have never placed an order. So, we print their names in the result set.

- **Approach 1:**

In this approach, we first try to find the customers who have ordered at least once. After having found this, we find the customers whose customer Id is not present in the previously obtained result set. This gives us the customers who have not placed a single order yet. The SQL query for it can be written as follows

```
select customers.name as 'Customers'
from customers
where customers.id not in
(
    select customerid from orders
);
```

- **Approach 2:**

In this approach, we use the concept of JOIN. We will LEFT JOIN the customer table with the order table based on the condition that id of the customer table must be equal to that of the customer id of the order table. Now, in our joined resultant table, we just need to find those customers whose order id is null. The SQL query for this can be written as follows:

```
select c.name as 'Customers' from Customers c
left join Orders o ON (o.customerId = c.id)
where o.id is null
```

Here, we first create aliases of the tables Customers and Orders with the name 'c' and 'o' respectively. Having done so, we join them with the condition that o.customerId = c.id. At last, we check for the customers whose o.id is null.

5. Given the following schema:

Table: Cinema

| Column Name | Type |
|-------------|---------|
| id | int |
| movie | varchar |
| description | varchar |
| rating | float |

The primary key for this table is id. Each row includes information about a movie's name, genre, and rating. rating is a float with two decimal digits in the range [0, 10]. Write an SQL query to report the movies with an odd-numbered ID and a description that is not "boring". Return the result table ordered by rating in descending order.

Example:

Input: Cinema table:

| id | movie | description | rating |
|----|---------|-------------|--------|
| 1 | War | thriller | 8.9 |
| 2 | Dhakkad | action | 2.1 |
| 3 | Gippi | boring | 1.2 |
| 4 | Dangal | wrestling | 8.6 |
| 5 | P.K. | Sci-Fi | 9.1 |

Output

| id | movie | description | rating |
|----|-------|-------------|--------|
| 5 | P.K. | Sci-Fi | 9.1 |
| 1 | War | thriller | 8.9 |

Explanation:

There are three odd-numbered ID movies: 1, 3, and 5. We don't include the movie with ID = 3 in the answer because it's boring. We put the movie with id 5 at the top since it has the highest rating of 9.1.

This question has a bit of ambiguity on purpose. You should ask the interviewer whether we need to check for the description to exactly match "boring" or we need to check if the word "boring" is present in the description. We have provided solutions for both cases.

- **Approach 1 (When the description should not be exactly "boring" but can include "boring" as a substring):**

In this approach, we use the MOD operator to check whether the id of a movie is odd or not. Now, for all the odd-numbered id movies, we check if its description is not boring. At last, we sort the resultant data according to the descending order of the movie rating. The SQL query for this can be written as follows:

```
select *
from cinema
where mod(id, 2) = 1 and description != 'boring'
order by rating DESC;
```

- **Approach 2 (When the description should not even contain "boring" as a substring in our resultant answer):**

In this approach, we use the LIKE operator to match the description having "boring" as a substring. We then use the NOT operator to eliminate all those results. For the odd-numbered id, we check it similarly as done in the previous approach. Finally, we order the result set according to the descending order of the movie rating. The SQL query for it can be written as follows:

```
SELECT *
FROM Cinema
WHERE id % 2 = 1 AND description NOT LIKE '%boring%'
ORDER BY rating DESC;
```

6. Consider the following schema:

Table: Users

| Column Name | Type |
|----------------|---------|
| account_number | int |
| name | varchar |

The account is the primary key for this table. Each row of this table contains the account number of each user in the bank. There will be no two users having the same name in the table.

Table: Transactions

| Column Name | Type |
|----------------|------|
| trans_id | int |
| account_number | int |
| amount | int |
| transacted_on | date |

trans_id is the primary key for this table. Each row of this table contains all changes made to all accounts. The amount is positive if the user received money and negative if they transferred money. All accounts start with a balance of 0.

Construct a SQL query to display the names and balances of people who have a balance greater than \$10,000. The balance of an account is equal to the sum of the amounts of all transactions involving that account. You can return the result table in any order.

Example:

Input: Users table:

| Account_number | name |
|----------------|-------|
| 12300001 | Ram |
| 12300002 | Tim |
| 12300003 | Shyam |

Transactions table:

| trans_id | account_number | amount | transacted_on |
|----------|----------------|--------|---------------|
| 1 | 12300001 | 8000 | 2022-03-01 |

| trans_id | account_number | amount | transacted_on |
|----------|----------------|--------|---------------|
| 2 | 12300001 | 8000 | 2022-03-01 |
| 3 | 12300001 | -3000 | 2022-03-02 |
| 4 | 12300002 | 4000 | 2022-03-12 |
| 5 | 12300003 | 7000 | 2022-02-07 |
| 6 | 12300003 | 7000 | 2022-03-07 |
| 7 | 12300003 | -4000 | 2022-03-11 |

Output:

| name | balance |
|------|---------|
| Ram | 13000 |

Explanation:

- Ram's balance is $(8000 + 8000 - 3000) = 11000$.
- Tim's balance is 4000.
- Shyam's balance is $(7000 + 7000 - 4000) = 10000$.
- **Approach 1:**

In this approach, we first create aliases of the given two tables' users and transactions. We can natural join the two tables and then group them by their account number. Next, we use the SUM operator to find the balance of each of the accounts after all the transactions have been processed. The SQL query for this can be written as follows:

SELECT u.name, SUM(t.amount) AS balance

FROM Users natural join Transactions t

GROUP BY t.account_number

HAVING balance > 10000;

7. Given the following schema:

Table: Employee

| Column Name | Type |
|-------------|----------|
| id | int |
| name | varcahar |
| department | varchar |
| managerId | int |

All employees, including their managers, are present at the Employee table. There is an Id for each employee, as well as a column for the manager's Id. Write a SQL query that detects managers with at least 5 direct reports from the Employee table.

Example:

Input:

| Id | Name | Department | ManagerId |
|-----|-----------|------------|-----------|
| 201 | Ram | A | null |
| 202 | Naresh | A | 201 |
| 203 | Krishna | A | 201 |
| 204 | Vaibhav | A | 201 |
| 205 | Jainender | A | 201 |
| 206 | Sid | B | 201 |

Output:

| Name |
|------|
| Ram |

- **Approach:**

In this problem, we first find all the manager ids who have more than 5 employees under them. Next, we find all the employees having the manager id present in the previously obtained manager id set.

The SQL query for this can be written as follows:

SELECT Name

FROM Employee

WHERE id IN

```
(SELECT ManagerId  
FROM Employee  
GROUP BY ManagerId  
HAVING COUNT(DISTINCT Id) >= 5);
```

Start Your Coding Journey With Tracks

Master Data Structures and Algorithms

Topic Buckets

Mock Assessments

Reading Material

8. Consider the following table schema:

Construct an SQL query to retrieve duplicate records from the Employee table.

Table: Employee

| Column Name | Type |
|-------------|---------|
| id | int |
| fname | varchar |
| lname | varchar |
| department | varchar |
| projectId | varchar |
| address | varchar |
| dateofbirth | varchar |
| gender | varchar |

Table: Salary

| Column Name | Type |
|---------------|---------|
| id | int |
| position | varchar |
| dateofJoining | varchar |
| salary | varchar |

Now answer the following questions:

1. Construct an SQL query that retrieves the fname in upper case from the Employee table and uses the ALIAS name as the EmployeeName in the result.

`SELECT UPPER(fname) AS EmployeeName FROM Employee;`

2. Construct an SQL query to find out how many people work in the "HR" department

`SELECT COUNT(*) FROM Employee WHERE department = 'HR';`

3. Construct an SQL query to retrieve the first four characters of the 'Iname' column from the Employee table.

`SELECT SUBSTRING(Iname, 1, 4) FROM Employee;`

4. Construct a new table with data and structure that are copied from the existing table 'Employee' by writing a query. The name of the new table should be 'SampleTable'.

`SELECT * INTO SampleTable FROM Employee WHERE 1 = 0`

5. Construct an SQL query to find the names of employees whose first names start with "S".

`SELECT * FROM Employee WHERE fname LIKE 'S%';`

6. Construct an SQL query to count the number of employees grouped by gender whose dateOfBirth is between 01/03/1975 and 31/12/1976.

`SELECT COUNT(*), gender FROM Employee WHERE dateOfBirth BETWEEN '01/03/1975' AND '31/12/1976' GROUP BY gender;`

7. Construct an SQL query to retrieve all employees who are also managers.

```
SELECT emp.fname, emp.Iname, sal.position
FROM Employee emp INNER JOIN Salary sal ON
emp.id = sal.id AND sal.position IN ('Manager');
```

8. Construct an SQL query to retrieve the employee count broken down by department and ordered by department count in ascending manner.

```
SELECT department, COUNT(id) AS DepartmentCount
FROM Employee GROUP BY department
ORDER BY DepartmentCount ASC;
```

9. Construct an SQL query to retrieve duplicate records from the Employee table.

```
SELECT id, fname, department, COUNT(*) as Count
FROM Employee GROUP BY id, fname, department
HAVING COUNT(*) > 1;
```

SQL Query Interview Questions for Experienced

1. Consider the following Schema:

Table: Tree

| Column Name | Type |
|-------------|------|
| id | int |
| parent_id | int |

Here, **id** is the primary key column for this table. **id** represents the unique identity of a tree node and **parent_id** represents the unique identity of the parent of the current tree node. The **id** of a node and the **id** of its parent node in a tree are both listed in each row of this table. There is always a valid tree in the given structure.

Every node in the given tree can be categorized into one of the following types:

1. "Leaf": When the tree node is a leaf node, we label it as "Leaf"
2. "Root": When the tree node is a root node, we label it as "Root"
3. "Inner": When the tree node is an inner node, we label it as "Inner"

Write a SQL query to find and return the type of each of the nodes in the given tree. You can return the result in any order.

Example:

Input: Tree Table

| id | parent_id |
|-----------|------------------|
| 1 | null |
| 2 | 1 |
| 3 | 1 |
| 4 | 3 |
| 5 | 2 |

Output:

| id | type |
|-----------|-------------|
| 1 | Root |
| 2 | Inner |
| 3 | Inner |
| 4 | Leaf |
| 5 | Leaf |

Explanation:

1. Because node 1's parent node is null, and it has child nodes 2 and 3, Node 1 is the root node.
2. Because node 2 and node 3 have parent node 1 and child nodes 5 and 4 respectively, Node 2 and node 3 are inner nodes.
3. Because nodes 4 and 5 have parent nodes but no child nodes, nodes 4, and 5 are leaf nodes.

Approach 1:

In this approach, we subdivide our problem of categorizing the type of each of the nodes in the tree. We first find all the root nodes and add them to our resultant set with the type "root". Then, we find all the leaf nodes and add them to our resultant set with the type "leaf". Similarly, we find all the inner nodes and add them to our resultant set with the type "inner". Now let us look at the query for finding each of the node types.

- **For root nodes:**

SELECT

id, 'Root' AS Type

FROM

tree

WHERE

parent_id IS NULL

Here, we check if the parent_id of the node is null, then we assign the type of node as 'Root' and include it in our result set.

- **For leaf nodes:**

SELECT

id, 'Leaf' AS Type

FROM

tree

WHERE

id NOT IN (SELECT DISTINCT

parent_id

FROM

tree

WHERE

parent_id IS NOT NULL)

AND parent_id IS NOT NULL

Here, we first find all the nodes that have a child node. Next, we check if the current node is present in the set of root nodes. If present, it cannot be a leaf node and we eliminate it from our answer set. We also check that the parent_id of the current node is not null. If both the conditions satisfy then we include it in our answer set.

- **For inner nodes:**

SELECT

id, 'Inner' AS Type

FROM

tree

WHERE

id IN (SELECT DISTINCT

parent_id

FROM

tree

WHERE

parent_id **IS NOT NULL**)

AND parent_id IS NOT NULL

Here, we first find all the nodes that have a child node. Next, we check if the current node is present in the set of root nodes. If not present, it cannot be an inner node and we eliminate it from our answer set. We also check that the parent_id of the current node is not null. If both the conditions satisfy then we include it in our answer set.

At last, we combine all three resultant sets using the UNION operator. So, the final SQL query is as follows:

SELECT

id, 'Root' **AS Type**

FROM

tree

WHERE

parent_id **IS NULL**

UNION

SELECT

id, 'Leaf' **AS Type**

FROM

tree

WHERE

```
id NOT IN (SELECT DISTINCT  
            parent_id  
           FROM  
           tree  
          WHERE  
            parent_id IS NOT NULL)  
          AND parent_id IS NOT NULL
```

UNION

```
SELECT  
    id, 'Inner' AS Type  
   FROM  
   tree  
  WHERE  
    id IN (SELECT DISTINCT  
            parent_id  
           FROM  
           tree  
          WHERE  
            parent_id IS NOT NULL)  
          AND parent_id IS NOT NULL  
  
ORDER BY id;
```

Approach 2:

In this approach, we use the control statement CASE. This simplifies our query a lot from the previous approach. We first check if a node falls into the category of “Root”. If the node does not satisfy the conditions of a root node, it implies that the node will either be a

“Leaf” node or an “Inner” node. Next, we check if the node falls into the category of “Inner” node. If it is not an “Inner” node, there is only one option left, which is the “Leaf” node.

The SQL query for this approach can be written as follows:

SELECT

id AS `Id`,

CASE

WHEN tree.id = (SELECT aliastree.id FROM tree aliastree WHERE aliastree.parent_id IS NULL)

THEN 'Root'

WHEN tree.id IN (SELECT aliastree.parent_id FROM tree aliastree)

THEN 'Inner'

ELSE 'Leaf'

END AS Type

FROM

 tree

ORDER BY `Id`;

Approach 3:

In this approach, we follow a similar logic as discussed in the previous approach. However, we will use the IF operator instead of the CASE operator. The SQL query for this approach can be written as follows:

SELECT

 aliastree.id,

 IF(ISNULL(aliastree.parent_id),

 'Root',

 IF(aliastree.id IN (SELECT parent_id FROM tree), 'Inner','Leaf')) Type

FROM

 tree aliastree

ORDER BY aliastree.id

2. Consider the following schema:

Table: Seat

| Column Name | type |
|-------------|---------|
| id | int |
| student | varchar |

The table contains a list of students. Every tuple in the table consists of a seat id along with the name of the student. You can assume that the given table is sorted according to the seat id and that the seat ids are in continuous increments. Now, the class teacher wants to swap the seat id for alternate students in order to give them a last-minute surprise before the examination. You need to write a query that swaps alternate students' seat id and returns the result. If the number of students is odd, you can leave the seat id for the last student as it is.

Example:

| id | student |
|----|---------|
| 1 | Ram |
| 2 | Shyam |
| 3 | Vaibhav |
| 4 | Govind |
| 5 | Krishna |

For the same input, the output is:

| id | student |
|----|---------|
| 1 | Shyam |
| 2 | Ram |
| 3 | Govind |
| 4 | Vaibhav |
| 5 | Krishna |

- Approach 1:

In this approach, first we count the total number of students. Having done so, we consider the case when the seat id is odd but is not equal to the total number of students. In this case, we simply increment the seat id by 1. Next, we consider the case when the seat id is odd but is equal to the total number of students. In this case, the seat id remains the same. At last, we consider the case when the seat id is even. In this case, we decrement the seat id by 1.

The SQL query for this approach can be written as follows:

SELECT

```
CASE WHEN MOD(id, 2) != 0 AND counts != id THEN id + 1 -- for odd ids
      WHEN MOD(id, 2) != 0 AND counts = id THEN id -- special case for last seat
      ELSE id - 1 -- For even ids
END as id,
```

student

FROM

```
seat, (SELECT COUNT(*) as counts
       FROM seat) AS seat_count
```

ORDER by id;

- **Approach 2:**

In this approach, we use the ROW_NUMBER operator. We increment the id for the odd-numbered ids by 1 and decrement the even-numbered ids by 1. We then sort the tuples, according to the id values. Next, we assign the row number as the id for the sorted tuples. The SQL query for this approach can be written as follows:

```
select row_number()
```

```
over (order by
```

```
(if(id%2=1,id+1,id-1))
```

```
) as id, student
```

```
from seat;
```

3. Given the following schema:

Table: Employee

| Column Name | type |
|--------------|---------|
| id | int |
| name | varchar |
| salary | int |
| departmentId | int |

id is the primary key column for this table. departmentId is a foreign key of the ID from the Department table. Each row of this table indicates the ID, name, and salary of an employee. It also contains the ID of their department.

Table: Department

| Column Name | type |
|-------------|---------|
| id | int |
| name | varchar |

id is the primary key column for this table. Each row of this table indicates the ID of a department and its name. The executives of an organization are interested in seeing who earns the most money in each department. A high earner in a department is someone who earns one of the department's top three unique salaries.

Construct a SQL query to identify the high-earning employees in each department. You can return the result table in any order.

Example:

Input: Employee table:

| id | name | salary | departmentId |
|----|-------|--------|--------------|
| 1 | Ram | 85000 | 1 |
| 2 | Divya | 80000 | 2 |
| 3 | Tim | 60000 | 2 |
| 4 | Kim | 90000 | 1 |
| 5 | Priya | 69000 | 1 |
| 6 | Saket | 85000 | 1 |

| id | name | salary | departmentId |
|-----------|-------------|---------------|---------------------|
| 7 | Will | 70000 | 1 |

Department table:

| id | name |
|-----------|-------------|
| 1 | Marketing |
| 2 | HR |

Output:

| Department | Employee | Salary |
|-------------------|-----------------|---------------|
| Marketing | Kim | 90000 |
| Marketing | Ram | 85000 |
| Marketing | Saket | 85000 |
| Marketing | Will | 70000 |
| HR | Divya | 80000 |
| HR | Tim | 60000 |

Explanation:

- Kim has the greatest unique income in the Marketing department - Ram and Saket have the second-highest unique salary.
- Will has the third-highest unique compensation.

In the HR department:

- Divya has the greatest unique income.
- Tim earns the second-highest salary.
- Because there are only two employees, there is no third-highest compensation.

Approach 1:

In this approach, let us first assume that all the employees are from the same department. So let us first figure out how we can find the top 3 high-earner employees. This can be done by the following SQL query:

```

select emp1.Name as 'Employee', emp1.Salary
from Employee emp1
where 3 >
(
  select count(distinct emp2.Salary)
from Employee emp2
where emp2.Salary > emp1.Salary
);

```

Here, we have created two aliases for the Employee table. For every tuple of the emp1 alias, we compare it with all the distinct salaries to find out how many salaries are less than it. If the number is less than 3, it falls into our answer set.

Next, we need to join the Employee table with the Department table in order to obtain the high-earner employees department-wise. For this, we run the following SQL command:

SELECT

 d.Name **AS** 'Department', e1.Name **AS** 'Employee', e1.Salary

FROM

 Employee e1

JOIN

 Department d **ON** e1.DepartmentId = d.Id

WHERE

 3 > (**SELECT**

COUNT(DISTINCT e2.Salary)

FROM

 Employee e2

WHERE

 e2.Salary > e1.Salary

AND e1.DepartmentId = e2.DepartmentId

);

Here, we join the Employee table and the Department table based on the department ids in both tables. Also, while finding out the high-earner employees for a specific department, we compare the department ids of the employees as well to ensure that they belong to the same department.

Approach 2:

In this approach, we use the concept of the DENSE_RANK function in SQL. We use the DENSE_RANK function and not the RANK function since we do not want the ranking number to be skipped. The SQL query using this approach can be written as follows:

```
SELECT Final.Department, Final.Employee, Final.Salary FROM
```

```
(SELECT D.name AS Department, E.name AS Employee, E.salary AS Salary,
```

```
DENSE_RANK() OVER (PARTITION BY D.name ORDER BY E.salary DESC) Rank
```

```
FROM Employee E, Department D
```

```
WHERE E.departmentId = D.id) Final
```

```
WHERE Final.Rank < 4;
```

Here, we first run a subquery where we partition the tuples by their department and rank them according to the decreasing order of the salaries of the employees. Next, we select those tuples from this set, whose rank is less than 4.

4. Given the following schema:

Table: Stadium

| Column Name | type |
|--------------|------|
| id | int |
| date_visited | date |
| count_people | int |

date_visited is the primary key for this table. The visit date, the stadium visit ID, and the total number of visitors are listed in each row of this table. No two rows will share the same visit date, and the dates get older as the id gets bigger. Construct a SQL query to display records that have three or more rows of consecutive ids and a total number of people higher than or equal to 100. Return the result table in ascending order by visit date.

Example:

Input: Stadium table:

| | id | date_visited | count_people |
|---|----|--------------|--------------|
| 1 | 1 | 2022-03-01 | 6 |
| 2 | 2 | 2022-03-02 | 102 |
| 3 | 3 | 2022-03-03 | 135 |
| 4 | 4 | 2022-03-04 | 90 |
| 5 | 5 | 2022-03-05 | 123 |
| 6 | 6 | 2022-03-06 | 115 |
| 7 | 7 | 2022-03-07 | 101 |
| 8 | 8 | 2022-03-09 | 235 |

Output:

| | id | date_visited | count_people |
|---|----|--------------|--------------|
| 5 | 5 | 2022-03-05 | 123 |
| 6 | 6 | 2022-03-06 | 115 |
| 7 | 7 | 2022-03-07 | 101 |
| 8 | 8 | 2022-03-09 | 235 |

Explanation:

The four rows with ids 5, 6, 7, and 8 have consecutive ids and each of them has ≥ 100 people attended. Note that row 8 was included even though the date_visited was not the next day after row 7.

The rows with ids 2 and 3 are not included because we need at least three consecutive ids.

- **Approach 1:**

In this approach, we first create three aliases of the given table and cross-join all of them. We filter the tuples such that the number of people in each of the alias' should be greater than or equal to 100.

The query for this would be

```
select distinct t1.*
```

```
from stadium t1, stadium t2, stadium t3
```

```
where t1.count_people >= 100 and t2.count_people >= 100 and t3.count_people >= 100;
```

Now, we have to check for the condition of consecutive 3 tuples. For this, we compare the ids of the three aliases to check if they form a possible triplet with consecutive ids. We do this by the following query:

```
select t1.*
```

```
from stadium t1, stadium t2, stadium t3
```

```
where t1.count_people >= 100 and t2.count_people >= 100 and t3.count_people >= 100
```

```
and
```

```
(
```

```
(t1.id - t2.id = 1 and t1.id - t3.id = 2 and t2.id - t3.id =1)
```

```
or
```

```
(t2.id - t1.id = 1 and t2.id - t3.id = 2 and t1.id - t3.id =1)
```

```
or
```

```
(t3.id - t2.id = 1 and t2.id - t1.id =1 and t3.id - t1.id = 2)
```

```
);
```

The above query may contain duplicate triplets. So we remove them by using the DISTINCT operator. The final query becomes as follows:

```
select distinct t1.*
```

```
from stadium t1, stadium t2, stadium t3
```

```
where t1.count_people >= 100 and t2.count_people >= 100 and t3.count_people >= 100
```

```
and
```

```
(
```

```
(t1.id - t2.id = 1 and t1.id - t3.id = 2 and t2.id - t3.id =1)
```

```
or
```

```
(t2.id - t1.id = 1 and t2.id - t3.id = 2 and t1.id - t3.id =1)
```

```
or
```

```
(t3.id - t2.id = 1 and t2.id - t1.id =1 and t3.id - t1.id = 2)
```

)

order by t1.id;

- **Approach 2:**

In this approach, we first filter out all the tuples where the number of people is greater than or equal to 100. Next, for every tuple, we check, if there exist 2 other tuples with ids such that the three ids when grouped together form a consecutive triplet. The SQL query for this approach can be written as follows:

with cte as

(select * from stadium

where count_people >= 100)

select cte.id, cte.date_visited, cte.count_people

from cte

where

((cte.id + 1) in (select id from cte)

and

(cte.id + 2) in (select id from cte))

or

((cte.id - 1) in (select id from cte)

and

(cte.id - 2) in (select id from cte))

or

((cte.id + 1) in (select id from cte)

and

(cte.id - 1) in (select id from cte))

Discover your path to a **Successful Tech Career!**

Answer 4 simple questions & get a career plan tailored for you

Interview Process

CTC & Designation

Projects on the Job

5. Given the following schema:

Table: Employee

| Column Name | Type |
|-------------|---------|
| id | int |
| company | varchar |
| salary | int |

Here, id is the id of the employee. company is the name of the company he/she is working in. salary is the salary of the employee Construct a SQL query to determine each company's median salary. If you can solve it without utilising any built-in SQL functions, you'll get bonus points.

Example:

Input:

| Id | Company | Salary |
|----|-----------|--------|
| 1 | Amazon | 1100 |
| 2 | Amazon | 312 |
| 3 | Amazon | 150 |
| 4 | Amazon | 1300 |
| 5 | Amazon | 414 |
| 6 | Amazon | 700 |
| 7 | Microsoft | 110 |

| Id | Company | Salary |
|-----------|----------------|---------------|
| 8 | Microsoft | 105 |
| 9 | Microsoft | 470 |
| 10 | Microsoft | 1500 |
| 11 | Microsoft | 1100 |
| 12 | Microsoft | 290 |
| 13 | Google | 2000 |
| 14 | Google | 2200 |
| 15 | Google | 2200 |
| 16 | Google | 2400 |
| 17 | Google | 1000 |

Output:

| Id | Company | Salary |
|-----------|----------------|---------------|
| 5 | Amazon | 414 |
| 6 | Amazon | 700 |
| 12 | Microsoft | 290 |
| 9 | Microsoft | 470 |
| 14 | Google | 2200 |